



TPU Beowulf Cluster

Friends of the TPU

Daan Dekoning Krekels

Bavo Debraekeleer

Jelte Boumans

Jesse Denaux

Tarun Singh

Inhoudstabel



• Doelstelling	<i>Tarun</i>
• Coral Dev Board	<i>Bavo</i>
• Mendel Linux	<i>Bavo</i>
• TensorFlow (Lite) model	<i>Tarun</i>
• Aanpak Cluster	<i>Daan & Tarun</i>
• Behuizing	<i>Jelte</i>
• Alternatieven	<i>Daan</i>
• Conclusie	<i>Daan</i>


A high-angle, isometric photograph of a Raspberry Pi 4 Model B single-board computer. The board is green and populated with various components. A large, silver, finned aluminum heat sink is mounted on the top, with a black 40mm cooling fan attached to its top surface. Two red wires are connected to the fan. The board features a USB-C port, a USB-A port, a 3.5mm audio jack, a micro-HDMI port, and a 40-pin GPIO header. A black SD card is inserted into the slot. The text "Doelstelling" is overlaid in white, with a red underline, across the center of the image.

Doelstelling

Doelstelling

- Cluster met TPU's bouwen
- Private netwerk
- Snelheid verhogen door de rekenkracht te verdelen
- NFS en Open MPI
- Runnen van een AI model op de cluster







Coral Dev Board



Single Board Computer

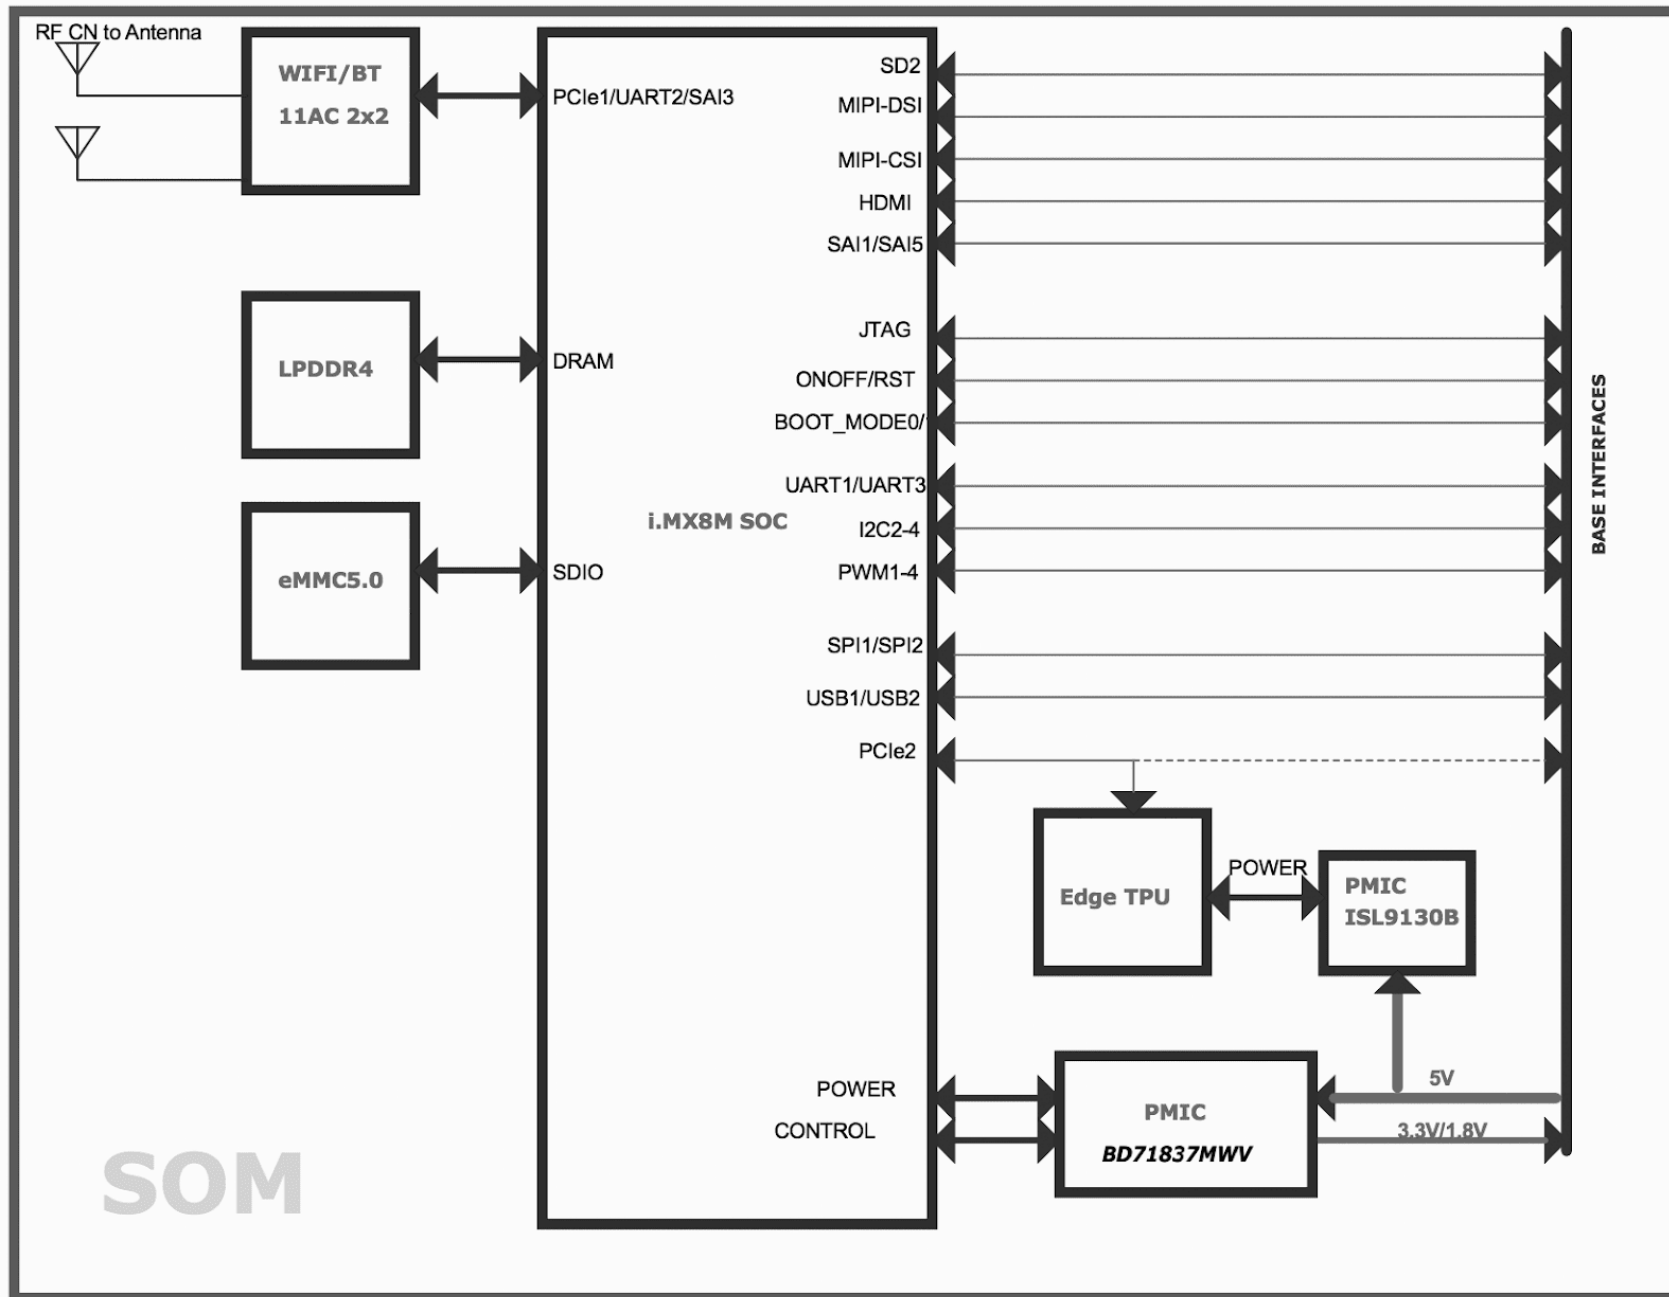
- Edge TPU System-on-Module (SoM)
- Baseboard met connectoren



Edge TPU SoM


- NXP i.MX 8M SoC (MIMX8MQ6DVA)
 - Cortex-A53/M4F
 - Vivante GC7000 Lite GPU en VPU
- Google Edge TPU ML accelerator coprocessor
- PMIC (BD71837MWV)
- Microchip ATECC608A Cryptographic coprocessor
- Wi-Fi 2x2 MIMO (802.11b/g/n/ac 2.4/5 GHz)
- Bluetooth 4.2
- 8/16 GB eMMC
- 1/4 GB LPDDR4

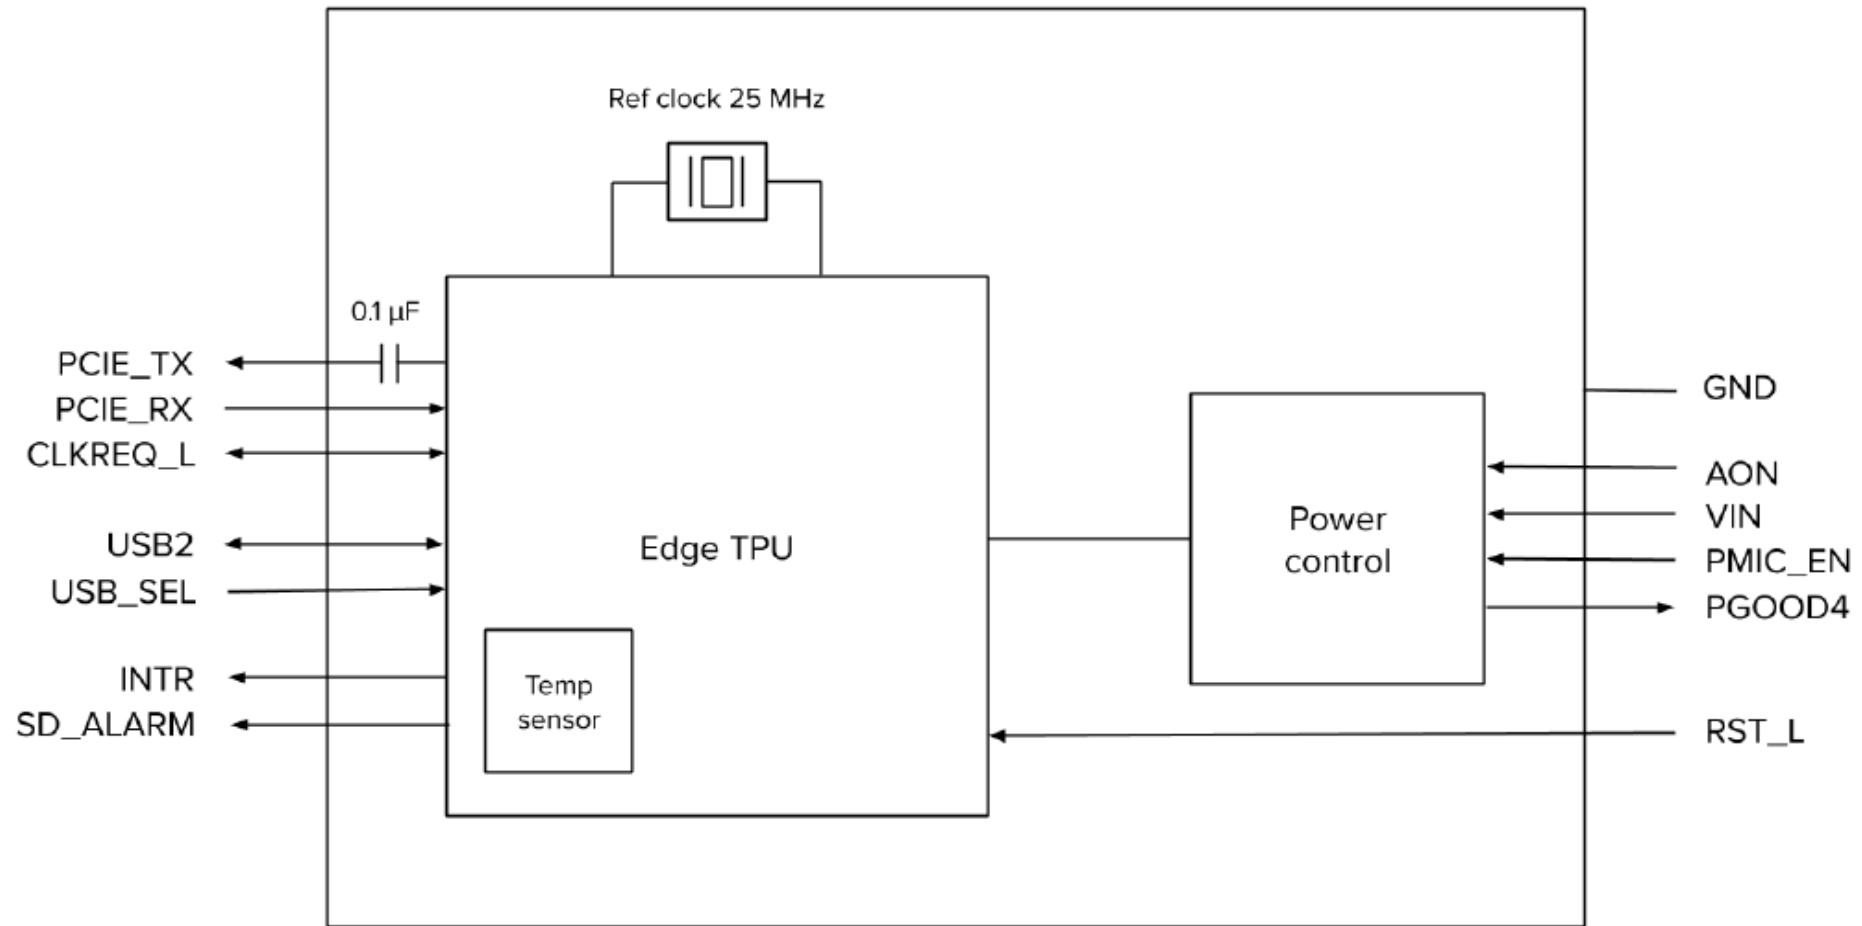




Google Edge TPU ML accelerator

- Tensor Processing Unit (TPU)
 - Google ASIC voor ML
 - Hoog volume laagprecisie-berekeningen
 - Laag energie verbruik
- Google Edge TPU ML accelerator
 - Lokaal TensorFlow Lite modellen
 - 4 TOPS @ 2 W
 - PMIC ISL9130B
 - PCIe Gen2 x1 interface







Google's TPUv3 has 8 cores — 4 chips each with 2 cores — with each core having 2 MXUs.



Baseboard

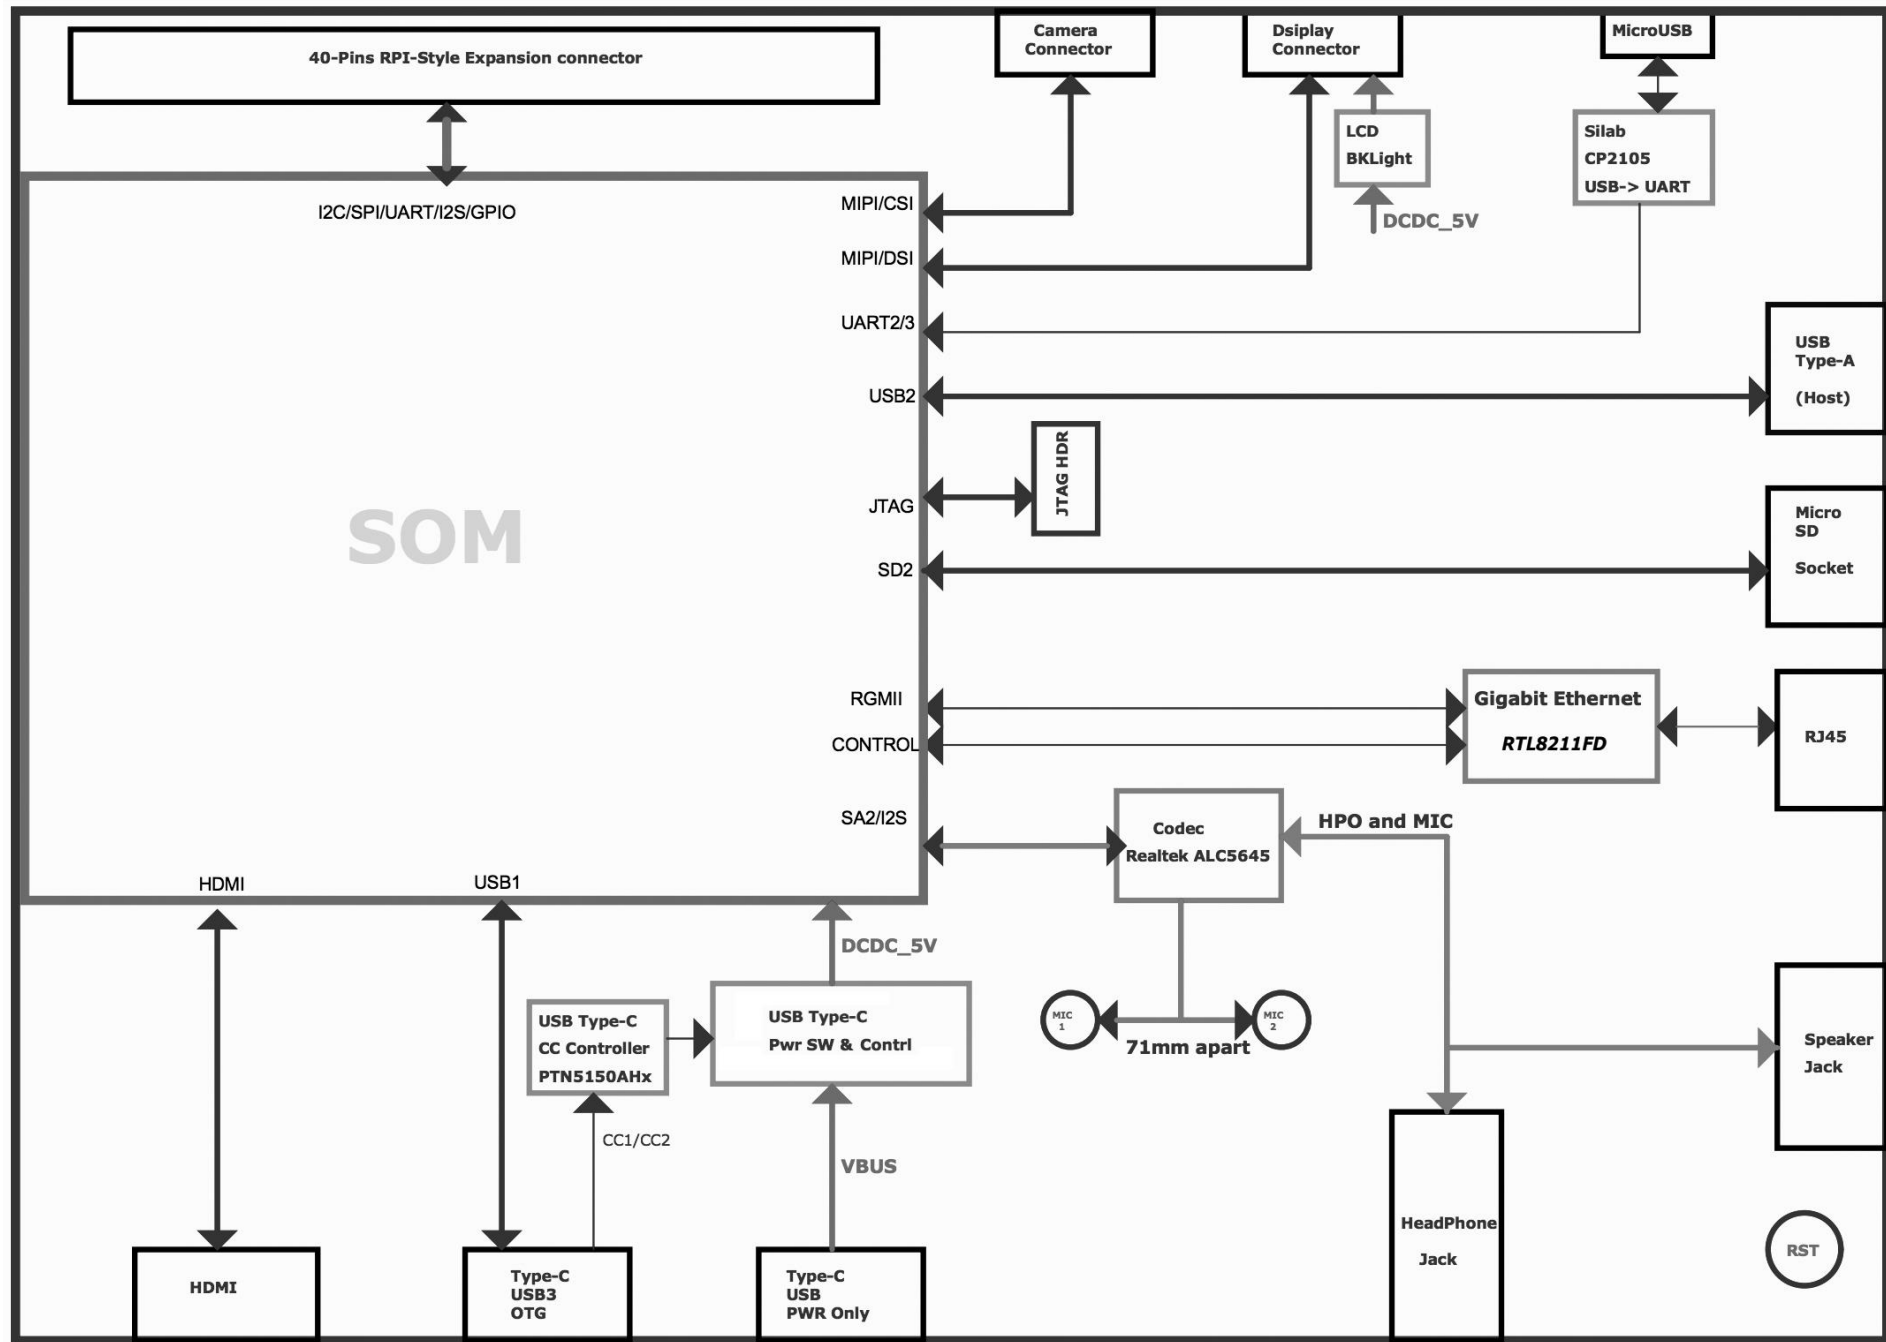
- Prototype board
- Connectiviteit


85 mm

56 mm











Mendel Linux

Wat is Mendel?

- *Lightweight derivative* upstream Debian 10 (buster) Linux distributie
- Voorgeïnstalleerde packages voor ML
 - Python 3.6 - 3.9
 - *python3-pycoral*
 - *libedgetpu1-std*
 - *python3-edgetpuvision*
 - *mendel-minimal*




Wat is Mendel?

- Aanpassing voor specifieke hardware Dev Board:
 - Device Tree Overlays
 - Coral "*apex*" PCIe driver
 - Edge TPU runtime
 - PyCoral library
 - Defconfig:
 - Include USB FTDI driver
 - Add micro k8s / k3s deps
 - i2s:
 - Enable I2S output on 40-pin header
 - rt5645:
 - Add a sysfs-based override for hp-detect

Model Pipelining

- Detecteerd PCI of USB peripherals


```
{'type': 'pci', 'path': '/dev/apex_0'}
```



```
enum edgetpu_device_type {  
    EDGETPU_APEX_PCI = 0,  
    EDGETPU_APEX_USB = 1,  
};  
  
struct edgetpu_device {  
    enum edgetpu_device_type type;  
    const char* path;  
};  
  
struct edgetpu_option {  
    const char* name;  
    const char* value;  
};
```

TPU Detectie

- `lspci -nn | grep 089a`
 - 0001:01:00.0 System peripheral [0880]:
Global Unichip Corp. Coral Edge TPU [1ac1:089a]
- `/sys/bus/pci/devices/0001:01:00.0`
- `/proc/bus/pci/devices`
 - 0100 1ac1089a 1eb 2020000c 0 2010000c 0 0 0 0 4000 0 100000 0 0 0 0 apex



Device Tree Overlay

- fsl-imx8mm-columbia.dts
 - *FSL i.MX8MM Columbia board*
 - Onderste laag, 2375 lijnen code
- fsl-imx8mq-phanbell.dts
 - *Freescale i.MX8MQ Phanbell*
 - Toevoeging "*google,edgetpu-audio-card*", *sound-rt5645*
 - Middelste laag, 2522 lijnen code
- fsl-imx8mq-yorktown.dts
 - Uitbreiding *Freescale i.MX8MQ Phanbell*
 - Toevoeging *Yorktown TPU "bus" device*, met 8x *Akira cores*, bestaande uit 2x *MXU's*
 - Bovenste laag, 2827 lijnen code



Device Tree Overlay

```
imx8mq-phanbell {  
  
    pcie1grp {  
        fsl,pins = < 0x230 0x498 0x528 0x12 0x00 0x76 0x13c 0x3a4 0x00 0x05 0x00 0x19 >;  
        linux,phandle = < 0x67 >;  
        phandle = < 0x67 >;  
    };  
};
```

```
pcie1: pcie@0x33c00000 {  
    compatible = "fsl,imx8mq-pcie", "snps,dw-pcie";  
    reg = <0x0 0x33c00000 0x0 0x400000>, <0x0 0x27f00000 0x0 0x80000>;  
    reg-names = "dbi", "config";  
    #address-cells = <3>;  
    #size-cells = <2>;  
    device_type = "pci";  
    ranges = <0x81000000 0 0x00000000 0x0 0x27f80000 0 0x00010000 /* do  
        0x82000000 0 0x20000000 0x0 0x20000000 0 0x07f00000>; /* non-  
    num-lanes = <1>;  
    interrupts = <GIC_SPI 74 IRQ_TYPE_LEVEL_HIGH>,  
        <GIC_SPI 80 IRQ_TYPE_LEVEL_HIGH>; /* eDMA */  
    interrupt-names = "msi";  
    #interrupt-cells = <1>;
```

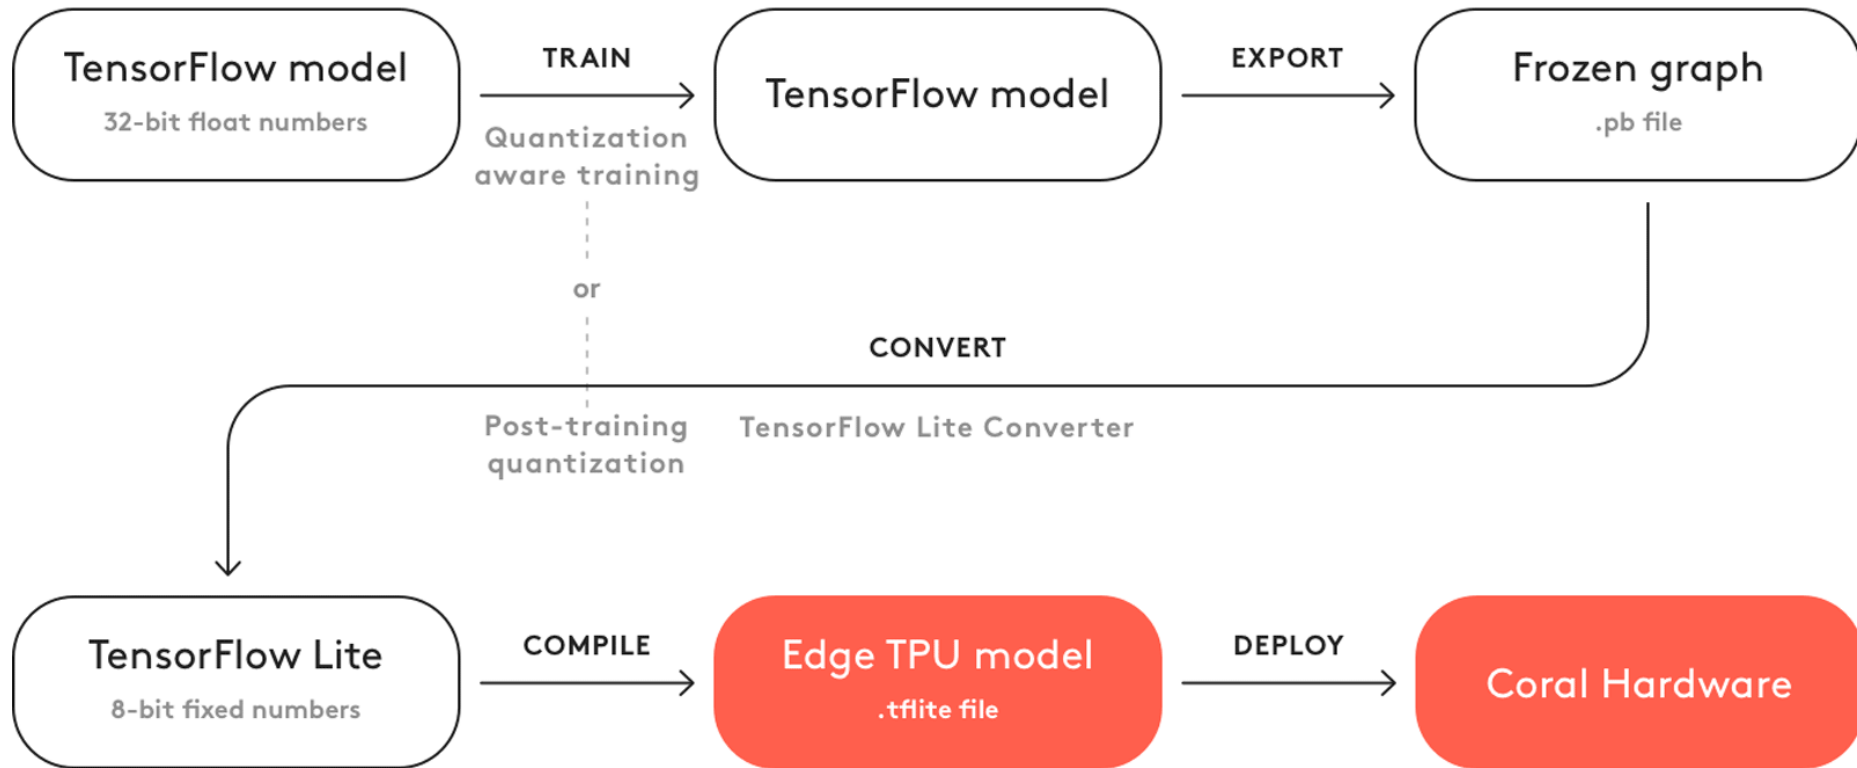
```
yorktown-tpu {  
    compatible = "google,yorktown-tpu";  
  
    akira0 {  
        compatible = "google,akira";  
        status = "okay";  
        tpu-gpios = < 0x6c 0x01 0x00 0x6c 0x00 0x01 >;  
    };  
  
    akira1 {  
        compatible = "google,akira";  
        status = "okay";  
        tpu-gpios = < 0x6d 0x01 0x00 0x6d 0x00 0x01 >;  
    };  
  
    akira2 {  
        compatible = "google,akira";  
        status = "okay";  
        tpu-gpios = < 0x6e 0x01 0x00 0x6e 0x00 0x01 >;  
    };  
  
    akira3 {  
        compatible = "google,akira";  
        status = "okay";  
        tpu-gpios = < 0x6f 0x01 0x00 0x6f 0x00 0x01 >;  
    };  
  
    akira4 {  
        compatible = "google,akira";  
        status = "okay";  
        tpu-gpios = < 0x70 0x01 0x00 0x70 0x00 0x01 >;  
    };  
  
    akira5 {  
        compatible = "google,akira";
```

A top-down view of a Raspberry Pi 4 single-board computer. It features a green PCB with various ports including USB-C, USB-A, and a micro-HDMI port. A large silver heat sink with a black cooling fan is mounted on the top. Two red wires are connected to the board. The text "TensorFlow (Lite) model" is overlaid in white, underlined with a red line.

TensorFlow (Lite) model

TensorFlow (Lite) model

- Tensorflow model omzetten naar .tflite bestand (8bit)
- .tflite model compileren






Aanpak Cluster

Algemene opbouw

- Mendel Linux
- Lokaal netwerk
- NFS (Ganesha)
- Open MPI
- Beheerders script

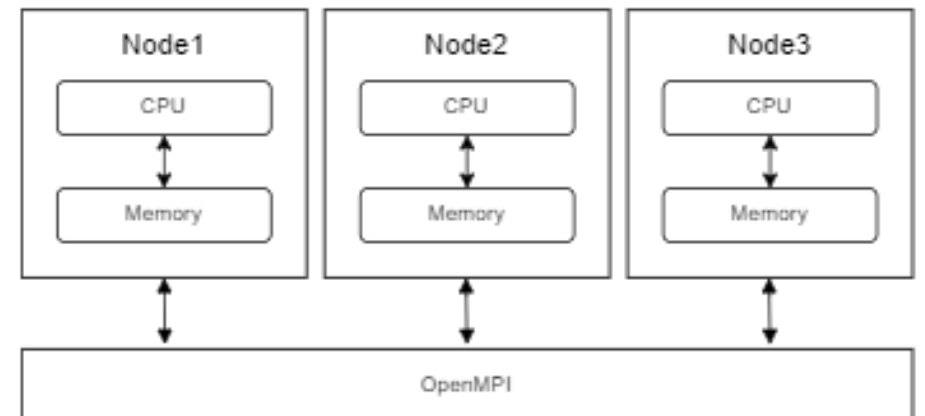




Network File System





OpenMPI

- Verdeling van taken
- Uitvoering van de taken door de nodes
- Resultaat terug naar de hoofd node sturen
- Gebruik van memory cores





Behuizing






Alternatieven

Alternatieven

- Gebruik van TensorFlow Serving
- Afstappen van Edge TPU Dev Board cluster



Alternatieven – TensorFlow Serving

Goals in Google datacenters

- Can serve multiple models, or multiple versions of the same model
- **Exposes both gRPC as well as HTTP inference endpoints**
- Allows deployment of new model versions without changing any client code
- Supports canarying new versions and A/B testing experimental models
- Adds minimal latency to inference time due to efficient, low-overhead implementation
- Features a scheduler that groups individual inference requests into batches for joint execution on GPU, with configurable latency controls
- Supports many servables: Tensorflow models, embeddings, vocabularies, feature transformations and even non-Tensorflow-based machine learning models

A top-down view of a Raspberry Pi single-board computer. It features a green PCB with various components including a USB port, a micro-USB port, a 3.5mm audio jack, a micro-HDMI port, and a 40-pin GPIO header. A large silver heat sink with a black cooling fan is mounted on the top. Two red wires are connected to the fan. The word "Conclusie" is overlaid in white text with a red underline.

Conclusie

Conclusie


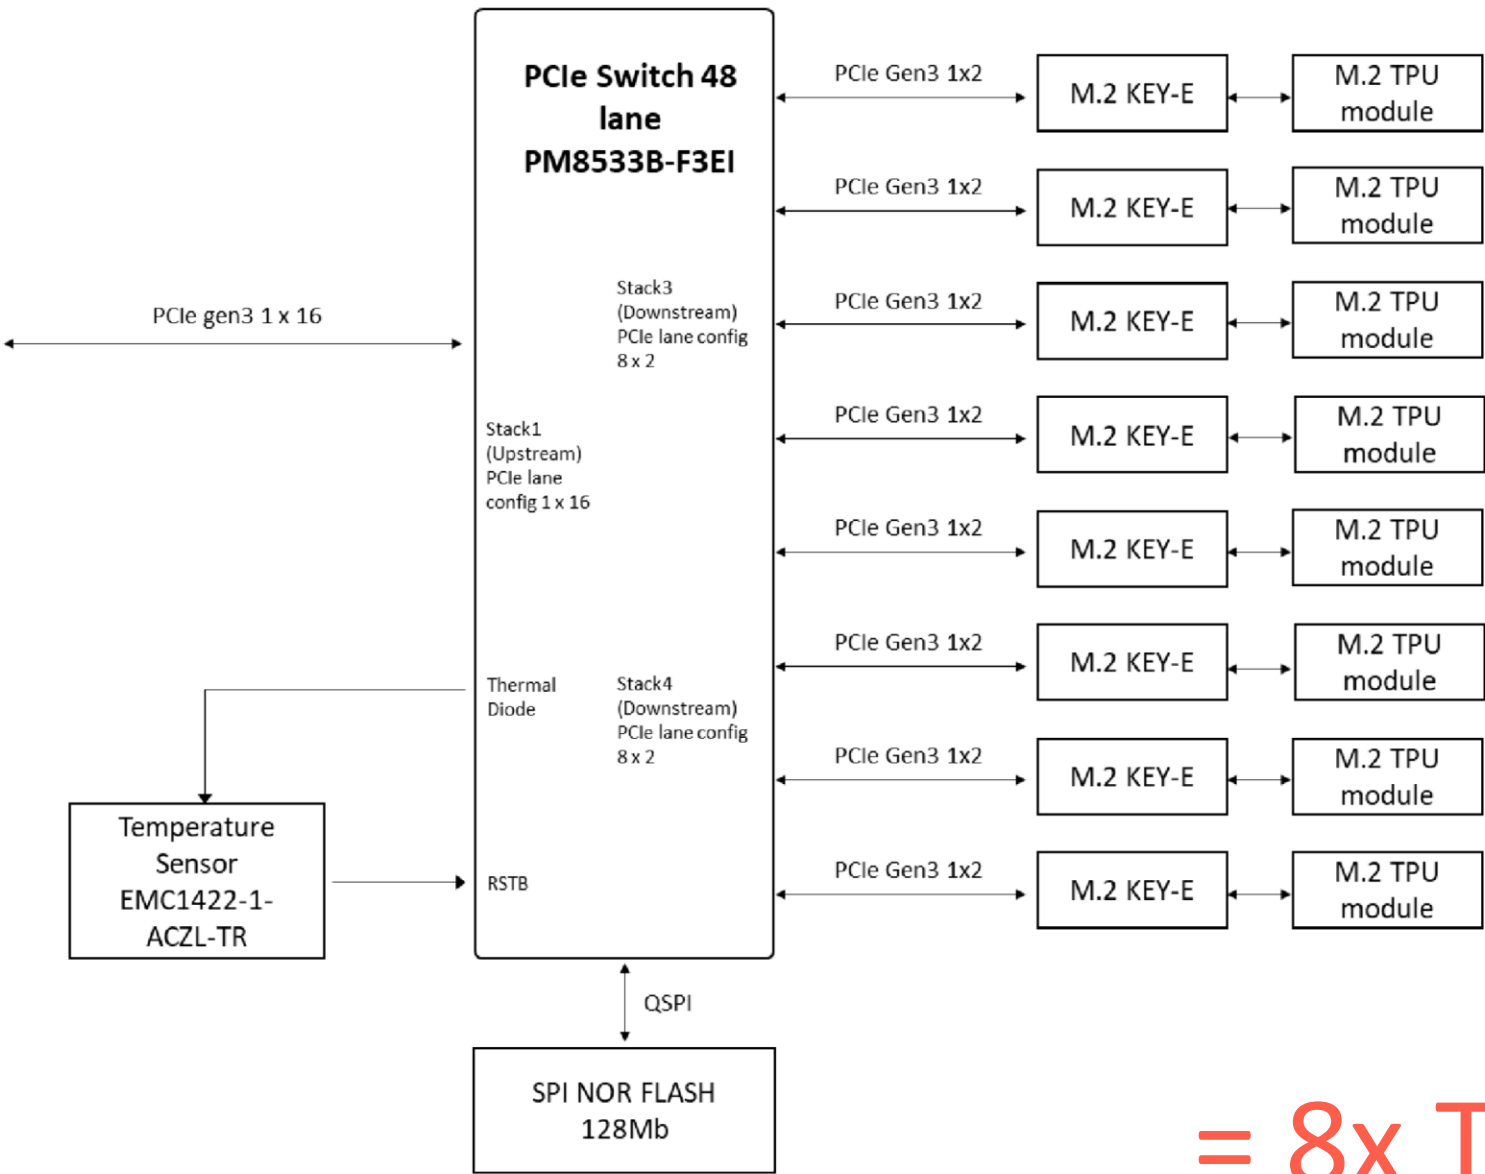
- Cluster van Coral Dev Boards moeilijk
- Vrijwel geen software ondersteuning
- Meerdere TPU modules op één "*normale*" computer
 - Goedkoper
 - Eenvoudiger
 - Bestaande software ondersteuning




PCIe x 16
Gold Finger

ATX POWER R/A
2x2 CONNECTOR
(spare)

FAN 4pin x 2
Header



= 8x TPU



Mini PCIe Accelerators



Biostar TB360-BTC PRO motherboard



= 11x TPU

